

TASK-4 (next)

Database using SQL

1. Database:

- A database is a collection of data.
- It is stored in such a way that we can store, manage, and access data easily.

2. SQL (Structured Query Language):

SQL is a language used to interact with databases.

Using SQL, we can Insert, Update, Delete, and Retrieve data.

3. Main SQL Commands:

O

CREATE DATABASE / TABLE → Create a new database or table.

• **INSERT INTO** → Add data.

O

SELECT → Retrieve/view data.

• **UPDATE** → Modify existing data.

• **DELETE** Remove

Keys in Database:

Primary Key → Unique identifier of a record.

• Foreign Key → Used to link two tables.

5. Joins:

• Used to combine data from multiple tables.

• Types: INNER JOIN, LEFT JOIN, RIGHT JOIN, FULL JOIN.

6. Functions in SQL:

O

Aggregate Functions: SUM(), AVG(), COUNT(), MIN(), MAX().

String Functions: CONCAT(), UPPER(), LOWER().

Date Functions: NOW(), DATE_ADD(), DATEDIFF().

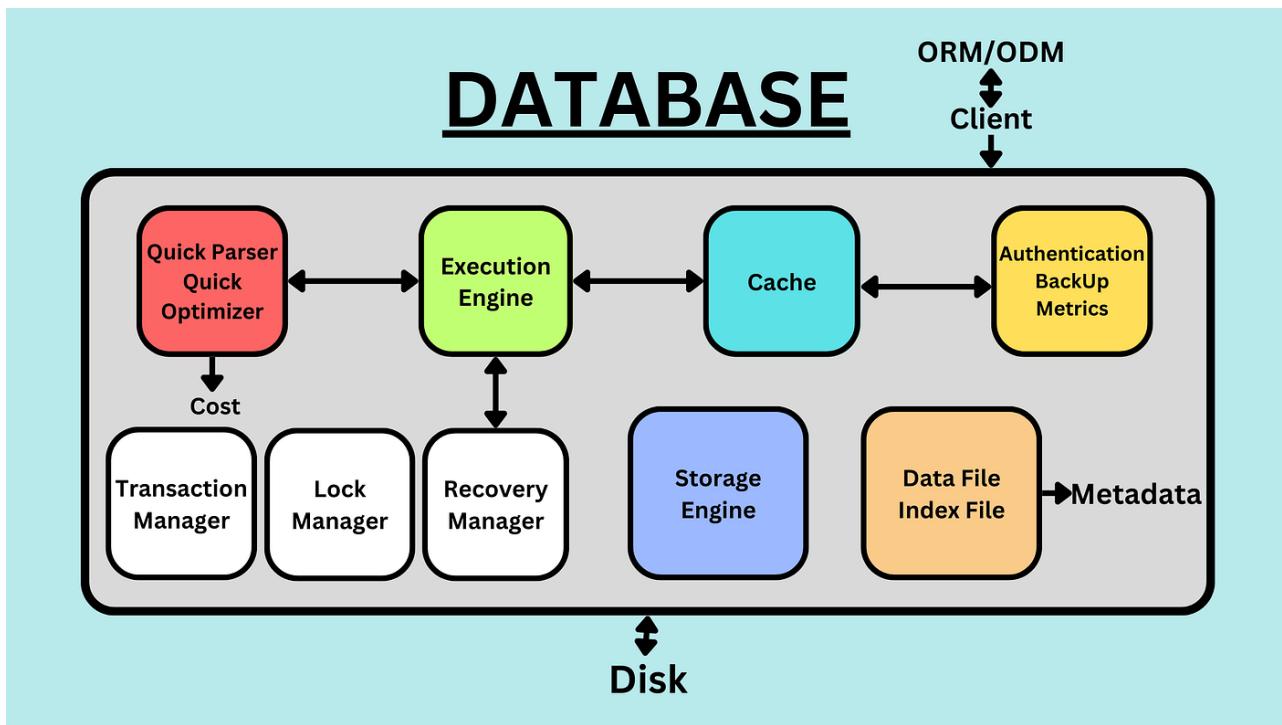
7. Views:

Save a query and use it like a virtual table.

8. Indexes:

O

Used to speed up data search



Database — deeper view

What it is: a structured store for persistent data (tables, indexes, constraints, transactions).

Types: relational (RDBMS: MySQL, PostgreSQL, SQLite), columnar (analytics), NewSQL/distributed, and NoSQL. Use RDBMS when relationships, transactions and strong consistency matter.

Modeling: start with an ER diagram (entities, relationships, attributes). Capture cardinality (1:1, 1:N, N:M) and optionality. Good modeling prevents slow queries later.

SQL (role & categories)

DDL (Data Definition Language):

CREATE

,

ALTER

,

DROP

— schema design, migration.

DML (Data Manipulation Language):

INSERT

,

UPDATE

,

DELETE

,

SELECT

— day-to-day data operations.

TCL (Transaction Control):

BEGIN/START

,

COMMIT

,

ROLLBACK

— atomic changes.

DCL (Data Control):

GRANT

,

REVOKE

— access control.

Best practice: Keep DDL changes in
version control (migration scripts).

Main SQL commands — practical examples & notes

CREATE TABLE (define types, constraints)

```
CREATE TABLE customers (
    customer_id SERIAL PRIMARY
KEY,
    name VARCHAR(200) NOT NULL,
    email VARCHAR(255) UNIQUE,
```

```
    created_at TIMESTAMP DEFAULT  
    NOW()  
);
```

INSERT

```
INSERT INTO customers (name,  
email) VALUES ('Asha',  
'asha@example.com');
```

SELECT

```
SELECT customer_id, name FROM  
customers WHERE email LIKE  
'%@example.com' ORDER BY  
created_at DESC LIMIT 10;
```

UPDATE / DELETE

```
UPDATE customers SET name =  
'A. Gupta' WHERE customer_id =  
5;
```

```
DELETE FROM customers WHERE  
created_at < '2020-01-01';
```

DROP: use with caution — usually via
migration scripts with backups.

Keys & constraints

Primary key: unique row id — prefer
narrow (small integer) for performance.

SERIAL

/

BIGSERIAL

or

UUID

.

Foreign key: enforces referential integrity;
can cascade
ON DELETE/UPDATE
or restrict.

Unique, NOT NULL, CHECK: ensure data
consistency at DB-level (defensive
programming).

Tradeoff: Constraints add CPU overhead
on writes but prevent data corruption and
simplify app logic.

Joins & subqueries — semantics & performance

INNER JOIN: returns rows matching both tables.

LEFT JOIN: keep all rows from left, matching or NULL from right.

RIGHT JOIN / FULL JOIN: use when appropriate; some RDBMS (MySQL) lack FULL JOIN.

Subqueries: correlated (run per outer row) vs non-correlated. Correlated can be slow; often rewrite as JOIN or use EXISTS

Example:

```
-- INNER JOIN
SELECT o.id, c.name, o.total
FROM orders o
```

```
JOIN customers c ON
o.customer_id = c.customer_id
WHERE o.total > 100;
```

```
-- Aggregation + HAVING
SELECT c.customer_id, COUNT(*) AS orders, SUM(o.total) AS total_spent
FROM customers c
```

```
JOIN orders o ON o.customer_id
= c.customer_id
GROUP BY c.customer_id
HAVING SUM(o.total) > 1000;
```

Performance tip: ensure join columns are indexed; avoid functions on join columns.

Functions & aggregates — use and optimization

Aggregates:

SUM, AVG, COUNT, MIN, MAX

. Use

GROUP BY

Window functions:

ROW_NUMBER(), RANK(), SUM()

OVER(PARTITION BY ...)

— powerful for running totals, top-N per group, without collapsing rows.

```
SELECT customer_id, order_id,
total,
    SUM(total) OVER
(PARTITION BY customer_id
ORDER BY order_date) AS
running_total
FROM orders;
```

Tip: grouping many rows is expensive—pre-aggregate if possible, use materialized views for frequently used aggregates.

Views & materialized views

Views: saved queries acting like virtual tables. Good for encapsulating complexity and implementing security (restrict columns).

Materialized views: store results physically — faster reads, but must be refreshed (manually or scheduled). Ideal for expensive aggregations.

Example:

```
CREATE VIEW active_customers AS

SELECT customer_id, name FROM
customers WHERE active = TRUE;

-- Materialized view (Postgres)
CREATE MATERIALIZED VIEW
monthly_sales AS

SELECT date_trunc('month',
order_date) AS month,
SUM(total) AS total
FROM orders GROUP BY 1;
```

Tradeoff: materialized views speed reads but add storage & refresh complexity.

Indexes — types & strategy

Common types: B-tree (default), Hash, GiST/GIN (Postgres: for full-text, arrays), BRIN (for very large ordered data).

Create index

```
CREATE INDEX
idx_orders_customer ON
orders(customer_id);
```

Composite index: order matters.

```
CREATE INDEX idx ON table(a,b);
```

helps queries filtering

a
or
a AND b
, but not only
b
.

Covering index: include columns so the query can be satisfied from the index alone
(

INCLUDE
in Postgres/MSSQL).

Partial/index on expressions: index only a subset
(
WHERE

clause) or expression (e.g., lower(email)).

Pitfalls: indexes speed reads, slow down writes. Too many indexes increases insert/update cost. Monitor index usage with

EXPLAIN
/statistics.

Selectivity matters: index columns with high distinct values (high cardinality).

Transactions & isolation

ACID: Atomicity, Consistency, Isolation, Durability.

Isolation levels: (from weak to strict)

READ UNCOMMITTED

,
READ COMMITTED

,
REPEATABLE READ

,
SERIALIZABLE

. Higher isolation reduces anomalies but increases locking/latency.

Example (Postgres style):

BEGIN;

```
UPDATE accounts SET balance =  
balance - 100 WHERE id = 1;
```

```
UPDATE accounts SET balance =  
balance + 100 WHERE id = 2;  
COMMIT;
```

Concurrency & locking: keep transactions short, avoid user input inside a transaction, use optimistic concurrency when possible (version columns), watch for deadlocks and use retry logic.

Schema design — normalization vs denormalization

Normalization: 1NF/2NF/3NF reduces redundancy and ensures integrity. Use it for OLTP workloads.

Denormalization: stores derived/duplicate data to speed reads (useful for read-heavy analytics). Adds complexity (need to keep duplicates in sync).

Partitioning & sharding: partition large tables by range/list/hash to speed queries and maintenance. Shard when dataset and throughput exceed single node capacity.

Query optimization checklist

- 1 Run
EXPLAIN
/
EXPLAIN ANALYZE
to see actual plan and cost. (EXPLAIN ANALYZE runs the query in Postgres).
- 2 Avoid
SELECT *
— choose only needed columns.

Use proper indexes (WHERE, JOIN, ORDER BY, GROUP BY columns).

- 4 Avoid functions on indexed columns (e.g., WHERE LOWER(name) = 'x' prevents index use unless you index the expression).
- 5 Use LIMIT + sensible pagination (cursor-based for large offsets).

Batch inserts/updates (bulk operations)

- 6 rather than many single-row operations.

Cache heavy, frequently-read results

- 7 (Redis, materialized views).

- 8 Maintain statistics (

ANALYZE

,

VACUUM

in Postgres).

Monitor slow query logs and optimize top

- 9 offenders.

Security & best practices

Least privilege: grant only required permissions to DB roles.

Use parameterized queries / prepared statements to prevent SQL injection.

Example (pseudo):

```
-- application code  
cursor.execute("SELECT * FROM users WHERE id = %s",  
(user_id,))
```

Encryption: encrypt connections (TLS), and sensitive data at-rest when required.

Audit & logging: enable logs for DDL, suspicious access, and monitor them.

Passwords & secrets: use secret managers, rotate credentials, use separate accounts for apps vs admins.

Scalability & high availability

Read replicas: scale reads by sending queries to replicas. Ensure replication lag is acceptable.

Master-master vs master-slave: choose based on write patterns. Master-master increases complexity for conflict resolution.

Sharding: split data horizontally when single node capacity is exceeded — application-level routing complexity.

Caching: use an in-memory cache (Redis / Memcached) to reduce DB load for frequent reads.

Connection pooling: essential for web apps (PgBouncer for Postgres, ProxySQL for MySQL).

Maintenance & monitoring

Backups: full & incremental, test restores regularly.

Maintenance tasks: vacuum/optimize, rebuild indexes, refresh materialized views during off-peak hours.

Monitoring metrics: QPS, latency, connection count, cache hit ratio, slow queries, index usage, replication lag. Use tools (Prometheus, Grafana, built-in RDBMS metrics).

Tools & workflows

Migrations: use tools (Flyway, Liquibase, Alembic) to manage schema changes.

CI/CD: run migrations in CI, test against staging DB.

ORMs: increase developer productivity but be aware of generated SQL; profile queries and use raw SQL for performance-sensitive paths.

Quick practical SQL examples (collected)

```
-- Create tables with PK + FK
```

```
CREATE TABLE products (
    product_id SERIAL PRIMARY KEY,
    name TEXT NOT NULL,
```

```
    price NUMERIC(12,2) NOT NULL
);
```

```
CREATE TABLE orders (
    order_id SERIAL PRIMARY KEY,
    customer_id INTEGER
REFERENCES
customers(customer_id),
    order_date TIMESTAMP DEFAULT
NOW(),
    total NUMERIC(12,2)
);
```

```
-- Join + aggregation
```

```
SELECT c.name,
COUNT(o.order_id) AS orders,
SUM(o.total) AS spent
FROM customers c
```

```
LEFT JOIN orders o ON
o.customer_id = c.customer_id
GROUP BY c.name
ORDER BY spent DESC;
```

```
-- Index for faster lookup
```

```
CREATE INDEX
idx_customers_email ON
customers (lower(email));

-- Transaction
BEGIN;

UPDATE accounts SET balance =
balance - 100 WHERE id = 1;

UPDATE accounts SET balance =
balance + 100 WHERE id = 2;
COMMIT;

-- EXPLAIN (Postgres)

EXPLAIN ANALYZE SELECT * FROM
orders WHERE customer_id = 123;
```

Practical checklist before you ship

Schema reviewed with constraints and indexes in place.

Critical queries run with

EXPLAIN ANALYZE

; slow queries optimized.

Backups configured and restore tested.

Monitoring and alerts for slow

queries/replication lag.

Least-privilege roles for applications.

Transactional boundaries and retry logic implemented in app.

If you want, I can:

analyze a specific schema or slow query you have (paste table schemas + query), or

produce optimized SQL examples for a sample ecommerce schema (customers, products, orders, order_items), or

A database is a collection of data.

It is stored in such a way that we can **store, manage, and access** data easily.

SQL is a language used to interact with databases.

Using SQL, we can **Insert, Update, Delete, and Retrieve** data.

CREATE DATABASE / TABLE → Create a new database or table.

INSERT INTO → Add data.

SELECT → Retrieve/view data.

UPDATE → Modify existing data.

DELETE → Remove data.

DROP → Delete a table or database completely.

Primary Key → Unique identifier of a record.

Foreign Key → Used to link two tables.

Used to combine data from multiple tables.

Types: INNER JOIN, LEFT JOIN, RIGHT JOIN, FULL JOIN.

Aggregate Functions: SUM(), AVG(),

COUNT(), MIN(), MAX().

String Functions: CONCAT(), UPPER(),

LOWER().

Date Functions: NOW(), DATE_ADD(),

DATEDIFF().

Save a query and use it like a virtual table.

Used to speed up data searching.

To ensure safe and controlled changes in a database.

Commands: COMMIT, ROLLBACK,
SAVEPOINT.